

Paperclip LLM Progress Brief

Current org chart, source links, operating intent, and evaluation context for an external LLM. Prepared May 20, 2026 from the live Paperclip API and the GitHub-backed Paperclip repo.

Evaluator Snapshot

Paperclip is Ian's control plane for AI-agent companies: agent registry, org structure, issues, approvals, costs, quota checks, worker routes, and durable operating records. The near-term intention is not more unattended autonomy. It is to make the system safe, observable, portable, cheap, and useful enough to support one real business pilot.

Live app	https://ians-mac-mini-1.tail403c1a.ts.net
Company	Ian Projects
Company id	a1a88815-804e-4282-ad9c-77107d763374
Host	Ian Mac Mini on Tailscale; repo path /Users/openclaw/Documents/paperclip
Portable repo	https://github.com/ianbel1914-netizen/paperclip.git
Working branch	codex/paperclip-control-plane-handoff
Current HEAD	b9d5a9b038510554a82caff676bcdea278a08e72
Important caveat	Windows checkout has uncommitted device-management files; this brief links stable committed docs and separately notes live API state.

Current Live Org Chart

Source: GET /api/companies/a1a88815-804e-4282-ad9c-77107d763374/org on May 20, 2026. Statuses are live operational posture, not the intended final design.

```
Ian / Board
  CEO [error]
    COO [error]
      Style [error]
        Chief of Staff [error]
          Runtime Watcher [error]
            Triage Agent [error]
              Reviewer Agent [error]
                Commit Agent [idle]
                  Workflow Fixer [error]
                    Cost Fixer [error]
                      Data Fixer [error]
                        Runtime Fixer [idle]
                          Fix Memory Agent [error]
                            Planning & Measurement [error]
                              Cost & Tooling Optimizer [error]
```

CTO [error]
 Coder [error]
 Coder 2 [error]
 Codex Engineer [idle]
 LLM TBD [paused]

Agent	Role / title	Status	Adapter	Reports to
CEO	ceo / Chief Executive Officer	error	claude_local	Board / Ian
COO	coo / Chief Operating Officer	error	claude_local	CEO
Style	style / Style Advisor	error	claude_local	COO
Chief of Staff	pm / Chief of Staff	error	claude_local	CEO
Runtime Watcher	runtime_watcher / Runtime Watcher	error	claude_local	Chief of Staff
Triage Agent	triage_agent / Triage Agent	error	claude_local	Chief of Staff
Reviewer Agent	fix_reviewer / Fix Reviewer	error	claude_local	Triage Agent
Commit Agent	commit_agent / Commit Agent	idle	codex_local	Reviewer Agent
Workflow Fixer	workflow_fixer / Workflow Fixer	error	claude_local	Triage Agent
Cost Fixer	cost_fixer / Cost Fixer	error	claude_local	Triage Agent
Data Fixer	data_fixer / Data Fixer	error	claude_local	Triage Agent
Runtime Fixer	runtime_fixer / Runtime Fixer	idle	codex_local	Triage Agent
Fix Memory Agent	fix_memory / Fix Memory Agent	error	claude_local	Triage Agent
Planning & Measurement	pm / Planning and Measurement Bot	error	claude_local	CEO
Cost & Tooling Optimizer	optimizer / Cost & Tooling Optimizer	error	claude_local	Planning & Measurement
CTO	cto / Chief Technology Officer	error	claude_local	CEO
Coder	engineer / Software Engineer	error	claude_local	CTO
Coder 2	engineer / Software Engineer	error	claude_local	CTO
Codex Engineer	engineer / Codex Software Engineer	idle	codex_local	CTO
LLM TBD	placeholder / LLM TBD Assignee	paused	claude_local	Board / Ian

What The Org Is Supposed To Do

- CEO: board-facing strategy, delegation, and escalation root. The human board/Ian remains the real authority.
- Chief of Staff Office: one front door for Ian, routing requests, clarifying outcomes, maintaining priorities, surfacing decisions, and protecting budget/attention.
- CTO and engineering lane: software implementation, repo work, testing, reviews, and controlled Codex/Claude worker routes.
- Planning & Measurement: outcome areas, worker-route metrics, quota visibility, budgets, and accountability evidence.

- Specialized fixer/reviewer lane: recovery, triage, cost, workflow, data, runtime, commit, and memory repair work under explicit control.
- LLM TBD: parked placeholder, not an active commitment.

Current Progress Signals

Signal	Current evidence	Why an evaluator should care
Portability	GitHub repo and handoff docs on branch codex/paperclip-control-plane-handoff	Lets work continue from Windows, Mac Mini, cloud Codex, or phone without relying on one chat.
Control gates	Quota Watch, routine dispatch guard, model fallback guard, stale-run handling, same-company quota endpoint access	Safety work is prioritized before broad autonomous execution.
Approvals	Approval workflow exists; recent intended UX is inline approve/reject/request-changes with plain-English Ian-question handling	Human approval remains a first-class governance primitive.
Mobile control	Noisy natural-language Telegram path was contained; a command-only Paperclip Control Bot was documented	The direction is deterministic remote control before conversational autonomy.
Portfolio model	Roadmap fronts and billing codes exist for Paperclip, CoS, EJV Revenue, IanOS, Real Estate, Acquisition, FamilyOS, Growth, Local Worker Pool, Memory	Progress is measured by outcome area rather than raw agent activity.
Live concern	Most live Claude-backed agents currently show error; Codex-backed agents queried show idle	The system is not ready for broad unattended work; remediation should start with runtime health and quota checks.

Intentions And Near-Term Sequence

1. Project 0: finish the Paperclip control-plane gate: safe, observable, low-cost agent work with green quota checks and smoke-test evidence.
2. Project 1: build the Chief of Staff operating layer as the one front door for Ian, initially as a design/operating pattern rather than an unrestricted new agent company.
3. Project 2: launch the EJV Labs Revenue Engine as the first business pilot after control gates pass.
4. Project 3: evolve IanOS and Real Estate Intelligence after one recurring value loop proves the operating model.
5. Parking lot: sports, ClickUp, and Notion are intentionally visible but not active execution fronts.

Repository Links For Review

- [Repo](#) - clickable GitHub source
- [Active branch](#) - clickable GitHub source
- [Current HEAD](#) - clickable GitHub source
- [Start-here handoff](#) - clickable GitHub source

- [Phone/Mac handoff](#) - clickable GitHub source
- [System memo](#) - clickable GitHub source
- [Project roadmap](#) - clickable GitHub source
- [Portfolio framework](#) - clickable GitHub source
- [Chief of Staff layer](#) - clickable GitHub source
- [Control bot](#) - clickable GitHub source
- [Model routing policy](#) - clickable GitHub source
- [Routine registry](#) - clickable GitHub source

Questions An LLM Should Answer

- Does the current architecture support the intended operating model without creating hidden autonomy or cost risk?
- Which live errors should be debugged first to restore a safe minimal team?
- Are the org chart and Chief of Staff routing rules too broad for the current control-plane maturity?
- Which next proof would most cheaply validate progress: quota recheck, Codex no-code smoke issue, approval workflow QA, or runtime health repair?
- Are GitHub docs, Paperclip issues, and live API state aligned, or are there drift points that need reconciliation?

Evaluation Boundaries

Do not infer that all agents should be resumed. Treat error statuses as a stop sign, not as permission to trigger recovery loops. Do not request or expose secrets. Prefer read-only live checks, GitHub review, narrow tests, and explicit Ian approval before waking agents, changing budgets, or enabling Telegram/Codex automation.

Source Notes

This brief used: live Paperclip API org/agent/dashboard reads, committed Git docs on the active branch, local git status/remotes/log, and prior Codex memory only to route which Paperclip fronts mattered. The live API result supersedes older handoff status notes where statuses differ.